

RISC vs CISC

In the Beginning...

- 1964 -- The first ISA appears on the IBM System 360
- In the “good” old days
 - Initially, the focus was on usability by humans.
 - Lots of “user-friendly” instructions (remember the x86 addressing modes).
 - Memory was expensive, so code-density mattered.
 - Many processors were *microcoded* -- each instruction actually triggered the execution of a builtin function in the CPU. Simple hardware to execute complex instructions (but CPIs are very, very high)
- ...SO...
 - Many, many different instructions, lots of bells and whistles
 - Variable-length instruction encoding to save space.
- ... their success had some downsides...
 - ISAs evolved organically.
 - They got messier, and more complex.


Things Changed

- In the modern era
 - Compilers write code, not humans.
 - Memory is cheap. Code density is unimportant.
 - Low CPI should be possible, but only for simple instructions
 - We learned a lot about how to design ISAs, how to let them evolve gracefully, etc.
- So, architects started with with a clean slate...

Reduced Instruction Set Computing (RISC)


- Simple, regular ISAs, mean simple CPUs, and simple CPUs can go fast.
 - Fast clocks.
 - Low CPI.
 - Simple ISAs will also mean more instruction (increasing IC), but the benefits should outweigh this.
- Compiler-friendly, not user-friendly.
 - Simple, regular ISAs, will be easy for compilers to use
 - A few, simple, flexible, fast operations that compiler can combine easily.
 - Separate memory access and data manipulation
 - Instructions access memory *or* manipulate register values. Not both.
 - “Load-store architectures” (like MIPS)

Instruction Formats



Arithmetic: $\text{Register}[\text{rd}] = \text{Register}[\text{rs}] + \text{Register}[\text{rt}]$

Register indirect jumps: $\text{PC} = \text{PC} + \text{Register}[\text{rs}]$



Arithmetic: $\text{Register}[\text{rd}] = \text{Register}[\text{rs}] + \text{Imm}$

Branches: If $\text{Register}[\text{rs}] == \text{Register}[\text{rt}]$, goto $\text{PC} + \text{Immediate}$

Memory: $\text{Memory}[\text{Register}[\text{rs}] + \text{Immediate}] = \text{Register}[\text{rt}]$

$\text{Register}[\text{rt}] = \text{Memory}[\text{Register}[\text{rs}] + \text{Immediate}]$



Direct jumps: $\text{PC} = \text{Address}$

Syscalls, break, etc.

RISC Characteristics of MIPS

- All instructions have
 - ≤ 1 arithmetic op
 - ≤ 1 memory access
 - ≤ 2 register reads
 - ≤ 1 register write
 - ≤ 1 branch
 - It needs a small, fixed amount of hardware.
- Instructions operate on memory *or* registers *not both*
 - “Load/Store Architecture”
- Decoding is easy
 - Uniform opcode location
 - Uniform register location
 - Always 4 bytes -> the location of the next PC is to know.
- Uniform execution algorithm
 - Fetch
 - Decode
 - Execute
 - Memory
 - Write Back
- Compiling is easy
 - No complex instructions to reason about
 - No special registers
- The HW is simple
 - A skilled undergrad can build one in 10 weeks.
 - 33 instructions can run complex programs.

CISC: x86

- x86 is the prime example of CISC (there were many others long ago)
 - Many, many instruction formats. Variable length.
 - Many complex rules about which register can be used when, and which addressing modes are valid where.
 - Very complex instructions
 - Combined memory/arithmetic.
 - Special-purpose registers.
 - Many, many instructions.
- Implementing x86 correctly is almost intractable

Mostly RISC: ARM

- ARM is somewhere in between
 - Four instruction formats. Fixed length.
 - General purpose registers (except the condition codes)
 - Moderately complex instructions, but they are still “regular” -- all instructions look more or less the same.
- ARM targeted embedded systems
 - Code density is important
 - Performance (and clock speed) is less critical
 - Both of these argue for more complex instructions.
 - But they can still be regular, easy to decode, and crafted to minimize hardware complexity
- Implementing an ARM processor is also tractable for 141L, but it would be harder than MIPS

RISCing the CISC

- Everyone believes that RISC ISAs are better for building fast processors.
- So, how do Intel and AMD build fast x86 processors?

- Despite using a CISC ISA, these processors are actually RISC processors. The preceding was a dramatization. MIPS instructions were used for clarity and because I had some laying around.

No x86 instruction were harmed in the production of this slide.

movb \$0x05, %al

movl -4(%ebp), %eax

movl %eax, -4(%ebp)

movl %R0, -4(%R1,%R2,4)

movl %R0, %R1

VLIWing the CISC

- We can also get rid of x86 in software.
- Transmeta did this.
 - They built a processor that was completely hidden behind a “soft” implementation of the x86 instruction set.
 - Their system would translate x86 instruction into an internal VLIW instruction set and execute that instead.
 - Originally, their aim was high performance.
 - That turned out to be hard, so they focused low power instead.
- Transmeta eventually lost to Intel
 - Once Intel decided it cared about power (in part because Transmeta made the case for low-power x86 processors), it started producing very efficient CPUs.

The End